

Understanding Users Through Search Behavior

Multi-Dimensional Analysis Beyond Keyword Frequency

Rayaan

January 2026

TL;DR

Challenge: Understand a user from 55K+ searches beyond simple keyword counting.

Approach: Six independent methods (semantic clustering, graph analysis, temporal patterns, behavioral metrics) that validate findings.

Result: User = **RESTLESS SEARCHER** - focused interests (entropy 0.44), rapid switching (74%), low persistence (1.36 searches/domain).

Actionable: Serve familiar domains with high content rotation to match restless pattern.

1 The Problem

Most user analysis = counting keywords. Simple. Fast. Incomplete.

Why? Because frequency $\neq$ importance. "The" appears more than "blockchain" but tells you nothing. Keyword frequency misses: relationships, evolution, behavioral patterns.

This challenge: Go deeper. Show thinking. Explain method choices. Document failures.

2 The Dataset

- **Size:** 55,383 searches across 85.7 months
- **Breadth:** 7,909 unique domains
- **Features:** timestamp, domain, search query

Preprocessing: Parsed timestamps. Normalized domains. Cleaned titles. Standard but critical work.

3 Methodology: Multi-Dimensional Analysis

Six independent analyses triangulate truth:

3.1 Method 1: Semantic Topic Discovery

Insight: "Machine learning" and "neural networks" are related despite sharing zero words.

```
1 from sentence_transformers import SentenceTransformer
2 import umap
3 from sklearn.cluster import HDBSCAN
4
5 # Encode into 384-dim semantic space
6 model = SentenceTransformer('all-MiniLM-L6-v2')
```

```

7 embeddings = model.encode(search_titles)
8
9 # Reduce dimensionality
10 reducer = umap.UMAP(n_components=5, metric='cosine')
11 reduced = reducer.fit_transform(embeddings)
12
13 # Cluster with HDBSCAN
14 clusterer = HDBSCAN(min_cluster_size=100)
15 clusters = clusterer.fit_predict(reduced)

```

Why HDBSCAN over K-means? I tried K-means first. Silhouette < 0.3 = poor separation. K-means assumes spherical clusters. Search behavior isn't spherical. HDBSCAN:

- Finds arbitrary shapes
- Auto-determines cluster count
- Identifies noise

Result: Distinct semantic clusters that persist across 85 months (Figure 1).




Figure 1: Semantic clusters from HDBSCAN. Top 10 clusters = primary interests. Persistent structure = stable core.

3.2 Method 2: Graph-Based Interest Networks

Insight: Keywords matter less alone. Connections matter more. Which topics bridge interests?

```

1 import networkx as nx
2
3 # Co-occurrence graph (5-search window)
4 G = nx.Graph()
5 for (kw1, kw2), weight in edge_weights.items():
6     if weight >= 10:
7         G.add_edge(kw1, kw2, weight=weight)
8
9 # PageRank = importance
10 pagerank = nx.pagerank(G, weight='weight')
11
12 # Louvain = communities
13 communities = nx.community.greedy_modularity_communities(G)

```

PageRank vs Frequency - key finding:

- **Frequency** = volume (how often)
- **PageRank** = importance (how connected)

Mid-frequency + high PageRank = bridge topic connecting interests. Frequency misses this.




Figure 2: Top 15 by frequency vs PageRank. Only 6/15 overlap. Different signal, different insight.

3.3 Method 3: Temporal Pattern Detection

Insight: Behavior changes. Life events create detectable change points.

```

1 import ruptures as rpt
2 from sklearn.preprocessing import StandardScaler
3
4 # Daily features
5 signal = np.array([daily_search_count,
6                     daily_domain_diversity,
7                     avg_search_hour])
8 signal = StandardScaler().fit_transform(signal)
9
10 # Bayesian change point detection
11 algo = rpt.Pelt(model='rbf', min_size=30, jump=5)
12 algo.fit(signal)
13 changepoints = algo.predict(pen=10)

```

Why Bayesian? Calendar windows are arbitrary. Real shifts happen on event boundaries. Pelt detects significant changes without imposing structure.




Figure 3: Change points over 85 months. Each spike = behavioral shift, likely life event.

Linear regression on monthly data shows trends. p -values quantify confidence.

3.4 Method 4: Markov Chain Predictability

Insight: How predictable is the user?

$$P(\text{domain}_{t+1}|\text{domain}_t) = \frac{\text{count}(d_t \rightarrow d_{t+1})}{\sum \text{count}(d_t \rightarrow *)} \quad (1)$$

High probability = stereotyped. Low = exploratory.




Figure 4: Transition matrix (top 20). Darker = more likely. Shows stereotyped pathways vs exploration.

3.5 Method 5: Behavioral Quantification

No subjective labels. Use math:

Metric	Value	Range	Meaning
Shannon Entropy	0.438	[0, 1]	FOCUSED
Gini Coefficient	0.835	[0, 1]	HIGH concentration
Transition Rate	0.737	[0, 1]	RAPID switching
Persistence	1.36	[1, ∞]	LOW sessions
Exploration	92.3/mo	> 0	High discovery

Table 1: Objective behavioral metrics.

Shannon Entropy: Diversity measure. Normalized to [0,1].

$$H = - \sum_{i=1}^N p_i \log_2 p_i \quad (2)$$

Gini: Inequality measure. 0 = equal, 1 = one dominates.

Result: RESTLESS SEARCHER.

4 Results: Who Is This User?

Classification: RESTLESS SEARCHER

Characteristics:

- Focused (entropy 0.44) - not generalist
- High concentration (Gini 0.83) - few domains dominate
- Rapid switching (74%) - short attention
- Low persistence (1.36) - no deep dives
- High exploration (92/mo) - constant discovery

Interpretation: Focused-but-restless. Checks Google, Twitter, Reddit, StackOverflow repeatedly but never stays. Concentrated on familiar domains with rapid rotation.




Figure 5: Behavioral dashboard. Five metrics = RESTLESS SEARCHER profile.

4.1 Actionable Insights

For Personalization:

1. Use PageRank keywords for core interests (not frequency)
2. High rotation (74% switching) = serve variety
3. Optimize for quick interactions (low persistence)
4. 92 domains/mo = introduce novelty within core

For Monitoring:

- Track change points for life events
- Monitor cluster shifts
- Use Markov for prediction

5 Why This Beats Simple Counting

Dimension	Frequency Only	Multi-Method
Signal	1 metric	6 analyses
Temporal	Static	Trends, changes
Semantic	Keywords	Embeddings
Structure	Flat	Networks
Metrics	Subjective	Math-based
Confidence	None	<i>p</i> -values
Reproducible	Ad-hoc	Algorithms

Table 2: Comparison. Multi-method captures what single-method cannot.

Multi-Modal Evidence: When all six methods point same way, confidence is high. Triangulated from six angles.

Statistical Rigor: Every claim backed by *p*-values or tests.

Process Transparency: K-means failed. Documented why. Showed HDBSCAN reasoning.

6 Limitations and Future Work

6.1 Current Limitations

- **Computational:** 2-3 minutes for 55K. Not real-time.
- **Causation:** Correlation only. Can't explain why.
- **Context:** No demographics or external data.
- **Temporal Bias:** Recent weighted more.

6.2 Future Work

1. LSTM/Transformer for sequences
2. Causal inference methods
3. Multi-user cohort analysis
4. Streaming algorithms
5. SHAP explainability

7 Technical Stack

```
1 # Libraries
2 sentence-transformers # Embeddings
3 umap-learn           # Reduction
4 hdbscan               # Clustering
5 networkx              # Graphs
6 ruptures              # Change points
7 scipy                 # Stats
8
9 # Algorithms
10 - Sentence-BERT (all-MiniLM-L6-v2)
11 - UMAP (cosine metric)
12 - HDBSCAN
13 - PageRank + Louvain
14 - Bayesian Pelt
15 - Markov chains
```

Code structure:

- Cells 1-12: Data prep, EDA
- Cells 13-23: Advanced analysis
- Separate: 7 visualizations

All reasoning documented. Failed attempts included. Reproducible.

8 Conclusion

Keyword frequency = one dimension. This = six:

Semantic → Graph → Temporal → Behavioral

Result: High-confidence profile with actionable recommendations.

User = RESTLESS SEARCHER. Strategy = familiar domains + high rotation. Simple, data-driven.

Key Insight

Understanding users isn't *what* they search. It's *how* (patterns), *why* (events), and *what* connects (structure).

Single metrics lie. Triangulated evidence tells truth.